

# AgenticAI Overview



Vaibhav Zodge

📞 7020616260

✉️ info.dnyanyog@gmail.com

🌐 <https://www.dnyanyog.org>

GitHub icon <https://github.com/zodgevaibhav>



# GenAI Program : Use Case



Gen AI is generation of Text, Image, Video or Audio using LLM

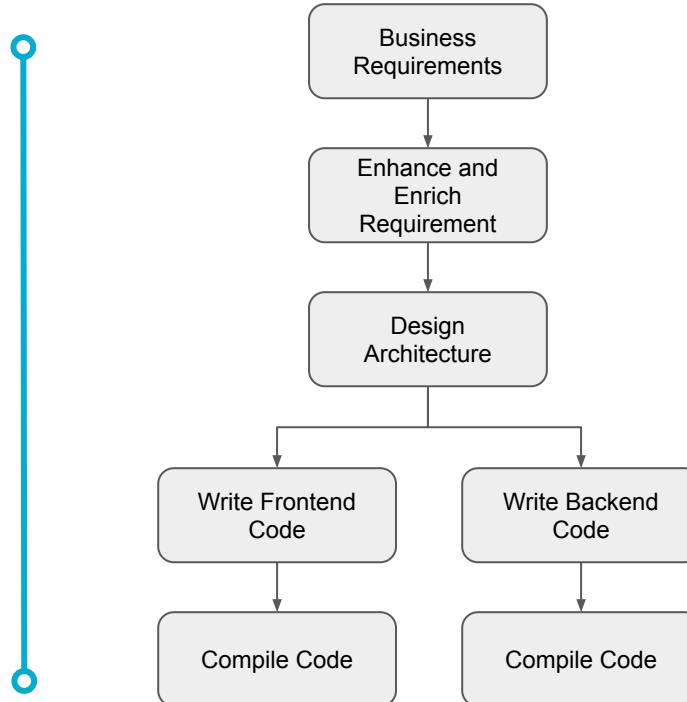
There are many real life usage of such system (GenAI)

## Agentic SDLC

- Read RAW Requirement
- Groom and document requirement
- Design Architect
- Write Backend and Frontend Code
- Compile and Build Code
- Write Unit Test
- Run Tests
- Fix Bugs
- Deploy application
- Monitor and improve

## Agentic Hospital Management

- Read the patient information
- Case History Summarization
- Help in diagnosis and prescription
- Appointment Schedule
- etc...



# AI SDLC



# Let's understand the Code



## Step 1: Requirement Refinement

- Raw requirements are given as input to the model
- A LLM is used to refine the raw requirements
- A predefined prompt guides the refinement process
- Output: **Well-structured and refined requirements**

```
from openai import OpenAI
import os, shutil
from datetime import datetime

client = OpenAI()

BASE = "runs"
FEATURES = """Feature: Patient Management API
Scenario: CRUD operations
"""

DEV_PROMPT = """
You are a Senior Full Stack Developer.

Generate a FULLY RUNNABLE project.

Backend:
- Spring Boot
- Maven
- REST CRUD
- Proper package structure
- Proper naming conventions
- Make sure to add CORS configuration to allow
  frontend-backend communication at localhost:3000
  port

Frontend:
- React
- npm
- Basic CRUD UI

STRICT RULES:
1. Use ONLY the format below
2. Every file MUST start with ===FILE===
3. Use full relative file paths
4. No explanations, no markdown
"""
```

```
def clean_workspace():
    if os.path.exists(BASE):
        shutil.rmtree(BASE)
    os.makedirs(BASE)

def run_agent(prompt, input_text):
    response = client.chat.completions.create(
        model="gpt-4o-mini",
        messages=[
            {"role": "system", "content": prompt},
            {"role": "user", "content": input_text}
        ],
        temperature=0.1
    )
    return response.choices[0].message.content

def write_code_files(output):
    blocks = output.split("==FILE==")
    for block in blocks:
        block = block.strip()
        if not block:
            continue
        lines = block.splitlines()
        path = lines[0]
        content = "\n".join(lines[1:])
        os.makedirs(os.path.dirname(path), exist_ok=True)
        with open(path, "w") as f:
            f.write(content)
    print(f"\n✓ {path}\n")
```

```
def main(requirement):
    clean_workspace()

    run_dir = os.path.join(BASE,
    f"run {datetime.now().strftime('%Y%m%d_%H%M%S')}")
    os.makedirs(run_dir)
    os.chdir(run_dir)

    # Requirement
    os.makedirs("requirements", exist_ok=True)
    with open("requirements/requirement.feature", "w") as f:
        f.write(requirement)

    # Architecture
    arch = run_agent(
        """you are a Software Architect.
        Produce architecture.md and PlantUML."""
        , requirement
    )
    os.makedirs("architecture", exist_ok=True)
    with open("architecture/architecture.md", "w") as f:
        f.write(arch)

    # Developer
    dev_output = run_agent(DEV_PROMPT, arch)
    materialize_files(dev_output)

    print("\n✓ FULL PROJECT GENERATED")

if __name__ == "__main__":
    main(FEATURES)
```

## Step 2: Architecture Design

- Refined requirements are used as input
- An LLM helps decide the application architecture
- A specific architecture prompt is used
- Output: **Architecture design document**

## Step 3: Code Generation

- Architecture document is given to the LLM
- LLM generates **Backend and Frontend code**
- Separate prompts are used for code generation
- Backend and Frontend generation functions run in parallel

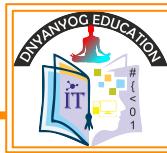
### What we needed ?

- Different Prompts and execution functions
- State Management (Output of one function as input to other)
- Decide the flow of execution (Requirement > Architecture > Backend Code > Frontend Code )
- Parallel Execution

### What More We want ?

- Retry mechanism (Iterative mechanism)
- Conditional LLM calls (routing & decision making)
- Observability & Debuggability
- .... & many more feature to make it work on production

# LangGraph



Lang Graph is a framework built on top of LangChain

Supports Graph Based workflows than Linear workflow

Graph like structure

Nodes → Steps/Agent/Function

Edges → Path that decide what runs next

Entry point → Decide which Node to start

End → Decide where to stop

Multiple Path

Conditional decision of next step

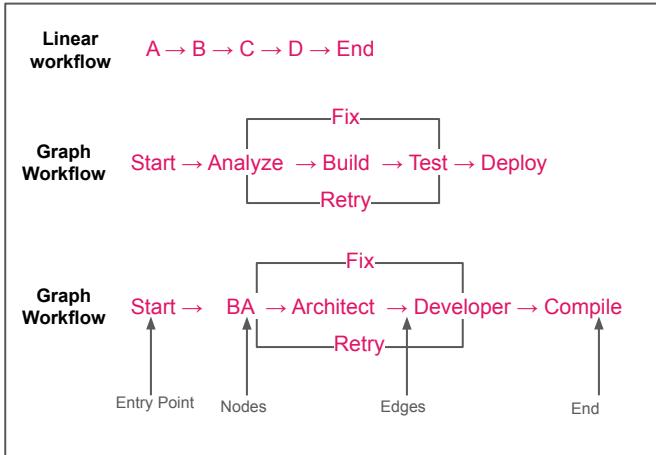
Loops possible

State Management between nodes/agents

Multiple agent management system

Conditional routing supported and inbuilt features for retries

Useful for multiple agent & decision based workflows, long running automation



Linear vs Graph Workflow



# Edges Examples

Sequential execution (Linear)

```
graph.add_edge("architect", "developer")
```



Parallel execution (Fan-out)

```
graph.add_edge("developer", "backend_compile")
graph.add_edge("developer", "frontend_compile")
```

Join execution (Fan-in / Sync)

```
graph.add_edge("backend_compile", "compile_done")
graph.add_edge("frontend_compile", "compile_done")
```

Loop execution (Retry / Iterate)

```
graph.add_edge("test", "fix_code")
graph.add_edge("fix_code", "test")
```

Conditional execution (Decision routing)

```
graph.add_conditional_edges(
    "test",
    lambda s: "pass" if s["ok"] else "fail",
    {"pass": "END", "fail": "fix_code"}
)
```



Early termination

```
graph.add_edge("deploy", END)
```

With **parallel execution** make sure shared keys are not mutated, else will get error  
Either use specific keys return  
ex.

**Wrong**

```
def backend_compiler_node(state):
    return state
```

```
def frontend_compiler_node(state):
    return state
```

**Correct**

```
def backend_compiler_node(state):
    return {"backend_build_ok": True}
```

```
def frontend_compiler_node(state):
    return {"frontend_build_ok": True}
```



# Monitoring using LangSmith

---



Create LangSmith Account

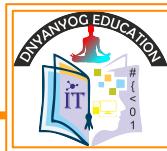
URL : <https://smith.langchain.com>

Use free tier and generate the api key

Set below parameters in environmental variables

```
export LANGSMITH_API_KEY= <<KEY>>
export LANGSMITH_TRACING=true
export LANGSMITH_ENDPOINT=https://api.smith.langchain.com
```

# Agentic Design Patterns



## Refinement Loop Pattern

### LangGraph primitives:

→ add\_edge() with a loop-back edge + quality check node

How: Routes output back to the same node until acceptance criteria are met.

## Retry with Feedback (Self-Healing) Pattern

### LangGraph primitives:

→ add\_conditional\_edges() + retry counter in state

How: On failure, routes back to the generator with error context.

## Guardrail / Gatekeeper Pattern

### LangGraph primitives:

→ Dedicated validation node + conditional routing

How: Validator node decides PASS → continue / FAIL → block or fix.

## Decompose–Solve–Compose Pattern

### LangGraph primitives:

→ Parallel nodes + merge/composer node

How: Subtasks write to different state keys and are combined later.

## Dual-Agent Verification Pattern

### LangGraph primitives:

→ Sequential nodes (Generator → Reviewer)

How: Reviewer updates approval or feedback in state.

## Conditional Routing Pattern

### LangGraph primitives:

→ add\_conditional\_edges()

How: Routing function chooses next node based on state values.

## Human-in-the-Loop Pattern

### LangGraph primitives:

→ Pause/wait node + external signal + resume edge

How: Workflow halts until human approval is received.

## Memory-Backed Decision Pattern

### LangGraph primitives:

→ External memory store + decision node

How: Node consults past outcomes to influence routing or prompts.

## Spec-Driven Generation Pattern

### LangGraph primitives:

→ Typed state (Pydantic / TypedDict) + schema validation node

How: Generator emits structured spec enforced by validator.

## Timeout & Escalation Pattern

### LangGraph primitives:

→ Retry counter in state + conditional exit edge

How: Exceeds threshold → escalate, fallback, or terminate safely.

## Sandbox Execution Pattern

### LangGraph primitives:

→ Execution node + error capture in state

How: Runs code/scripts in isolation and records logs/errors.

## Fallback Model / Tool Pattern

### LangGraph primitives:

→ Conditional routing + alternate agent/tool nodes

How: Failure routes to backup model or tool.

## Validation-First Pattern

### LangGraph primitives:

→ Entry validation node + early termination

How: Inputs are validated before costly generation steps.

## Checkpoint & Replay Pattern

### LangGraph primitives:

→ State checkpoint + resumable graph

How: Enables replay, debugging, and auditability.

## Progressive Disclosure Pattern

### LangGraph primitives:

→ Sequential nodes with incremental context enrichment

How: Each node adds detail before passing state forward.

**“Production agent systems succeed by structure, validation, and control — not by prompt size.”**

# Agentic AI : Note



Agentic systems are **state machines with intelligence**, not autonomous magic.

Reliability comes from **explicit state transitions**, not from clever prompts.

Every agent must have **one responsibility** and a clear success criterion.

If an agent cannot be **validated**, it should not be automated.

Good agent architectures **fail fast, fail safe, and recover automatically**.

Most “intelligence” in production agents lives in **routing logic**, not in LLMs.

Agent workflows should be **observable first, intelligent second**.

Parallelism increases speed but **multiplies state-management complexity**.

Determinism is achieved through **constraints, schemas, and retries**, not temperature.

A workflow without guardrails is a **liability, not an AI system**.

Agent loops must always have **explicit exit conditions**.

Memory is useful only when it **changes future decisions**.

Agents should produce **machine-consumable outputs**, not human prose.

Validation is cheaper than regeneration — **validate early, validate often**.

Production agents are **composed systems**, not single “smart” agents.

Human-in-the-loop is a **design choice**, not a failure mode.

Agent failures should **teach the system**, not just retry blindly.

Code execution is an agent too — **treat tools as first-class nodes**.

Good agent graphs are **boring, predictable, and repeatable**.

If a workflow cannot be diagrammed, it is probably **not production-ready**.

# Agentic AI : Note



## 1. Entry & Control Layer

Accepts raw input (user / system / event)  
Performs **early validation & sanitization**  
Initializes **shared workflow state**

*Fail fast before intelligence is applied*

## 2. Orchestration Layer (Graph Engine)

Deterministic **state machine**  
Nodes = agents / tools  
Edges = control flow & decisions  
Conditional routing & loops

*Control flow is the backbone of agentic systems*

## 3. Agent Layer (Single Responsibility)

Requirement Refiner  
Planner / Architect  
Generator (Code / Content / Spec)  
Reviewer / Critic  
Executor (Build / Test / Run)

*Agents think locally, architecture thinks globally*

## 4. State Layer (Source of Truth)

Typed, explicit, versioned state  
Immutable updates  
Retry counters, flags, errors

*State is the memory, context, and contract*

## 5. Validation & Guardrails

Schema validation  
Business rules  
Security & policy checks  
Output gating

*Nothing proceeds without validation*

## 6. Execution & Tooling

Sandboxed code execution  
CLI, Docker, APIs, scripts  
Logs & errors captured into state

*Tools are first-class citizens, not side effects*

## 7. Memory & Learning

Past failures & decisions  
Conventions & preferences  
Optional long-term storage

*Memory exists only to change future behavior*

## 8. Recovery & Escalation

Retry with feedback  
Fallback models/tools  
Human-in-the-loop  
Safe termination

*Self-healing beats perfect generation*

## 9. Observability & Persistence

State checkpoints  
Agent outputs  
Replay & audit

*If you can't replay it, you can't trust it*

*Agentic AI systems succeed by structure, validation, and control — not by prompt size.*

*LangGraph patterns emerge from how you combine nodes, state, and conditional edges — not from special features.*

